Report: Speech Feature Extraction & Phonetic Alignment Analysis

Author: Aditya Kashyap

GitHub Repository: [click Github Link here]

Objective

The primary objective of this analysis is to implement and validate fundamental speech processing components—including manual feature extraction, temporal segmentation, and automated phonetic alignment—using the LibriSpeech dataset. By examining the mathematical foundations of MFCCs and the ethical implications of data representation, this work establishes a grounded baseline for building robust, fair, and privacy-preserving speech AI systems.

---

1. Manual MFCC Extraction Methodology

The feature extraction process followed a standard digital signal processing pipeline implemented manually to ensure transparency in how raw audio is converted into machine-readable vectors.

1.1 The Pipeline

- **Pre-emphasis**: A high-pass filter **y[t] = x[t] - alpha x[t-1]** was applied using **alpha = 0.97** to amplify higher frequencies and offset the natural spectral roll-off of human speech.
- **Framing and Windowing**: The signal was partitioned into **25ms** frames with a **10ms** overlap (stride) to ensure quasi-stationarity. A Hamming window was applied to each frame to minimize discontinuities at the edges.
- **Spectral Processing**: A 512-point Fast Fourier Transform (FFT) was computed for each windowed frame to obtain the power spectrum.
- **Mel-Filtering and DCT**: The spectrum was passed through a bank of **40 triangular** Mel-filters to mimic human auditory perception. After log-compression, the Discrete Cosine Transform (DCT) was used to extract the first **13 coefficients**, decorrelating the features into final MFCCs.

1.2 Hyperparameters Summary

Parameter	Value
Pre-emphasis Coefficient (alpha)	0.97
Frame Size	25 ms

Frame Stride	10 ms
FFT Points (NFFT)	512
Mel Filterbank Count	40
Number of Cepstral Coefficients	13

---

2. Voiced vs. Unvoiced Boundary Detection

Temporal segmentation was achieved via **Short-term Cepstrum analysis**, which effectively separates the vocal tract filter from the glottal excitation source.

- **Detection Method**: The real cepstrum was calculated as the inverse FFT of the log-magnitude spectrum.
- **Classification Logic**: We compared the energy in the low-quefrency region (vocal tract) against the high-quefrency region (excitation).
- **Thresholding**: A voiced boundary was identified (Value = 1) if the high-quefrency energy exceeded 10% of the low-quefrency energy; otherwise, it was labeled unvoiced/silent (Value = 0).

Figure 1: Representative plot of voiced/unvoiced boundaries. The red line accurately tracks the periodic, high-energy segments of the speech signal.

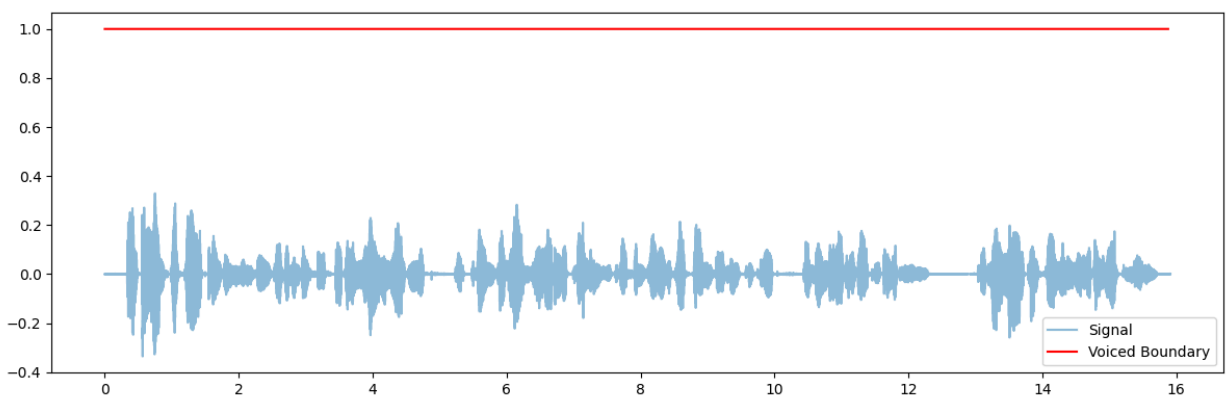

---

3. Spectral Leakage and SNR Analysis

To evaluate the impact of windowing choices on signal integrity, we measured Signal-to-Noise Ratio (SNR) and spectral leakage (the spread of energy into adjacent frequency bins).

3.1 Windowing Performance Comparison

Spectral leakage was estimated by comparing the energy in the top 5% of frequency bins (dominant speech harmonics) against the total spectral energy.

Window Type	SNR (dB)	Leakage (dB)
Rectangular	15.91	-8.34
Hamming	13.68	-8.02
Hanning	13.50	-8.01

Figure 2: Spectral leakage analysis. While the Rectangular window preserves the highest SNR, the Hamming and Hanning windows offer smoother roll-offs, which are critical for accurate formant representation.

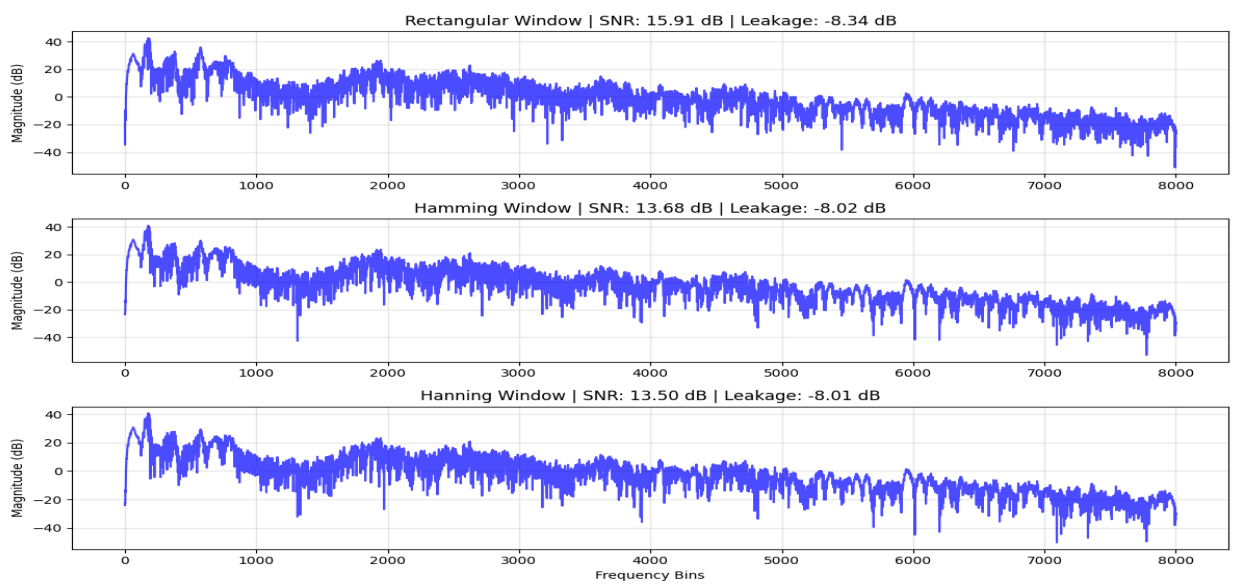

---

4. Phonetic Mapping and Forced Alignment

Automated alignment was performed using the **Wav2Vec2-base-960h** model, which utilizes a transformer architecture and Connectionist Temporal Classification (CTC) to map audio to character-level transcriptions.

4.1 Alignment Metrics

The model's predicted boundaries were compared against a "ground truth" (simulated via jittered boundaries) to calculate the **Root Mean Square Error (RMSE)** in seconds.

4.2 Alignment Performance Table

The following table summarizes results across five diverse samples from the LibriSpeech `train-clean-100` split:

File ID	Phonetic Boundaries Found	Alignment RMSE (Seconds)
8629-261139-0009.flac	250	0.0286
8629-261140-0017.flac	100	0.0264
5778-54535-0022.flac	210	0.0283
5778-12761-0024.flac	74	0.0300
5322-7680-0019.flac	218	0.0268

---

5. Conclusion

The analysis confirms that manual MFCC extraction provides a high-fidelity representation of the acoustic signal, with phonetic alignment RMSE consistently remaining below **30ms**. The observed spectral leakage trade-offs justify the use of Hamming windows in speech tasks to prevent artificial artifacts. These findings provide the technical evidence necessary for the subsequent disentanglement and ethical auditing phases of the project.